

# Developing and testing a device driver without the device

```
git clone --recurse-submodules git@github.com:mdionisio/linuxlab2018.git
```

# BEFORE WE START

```
git clone --recurse-submodules git@github.com:mdionisio/linuxlab2018.git

apt-get install build-essential
apt-get install crossbuild-essential-armhf
apt-get install gcc-arm-linux-gnueabi
apt-get install gcc-arm-none-eabi
apt-get install libncurses5-dev
apt-get install bison flex libssl-dev bc
apt-get install python pkg-config libgl2.0-dev libpixmap-1-dev
apt-get install sudo cpio lzop gzip
apt-get install gdb-multiarch
```

# Who we are

# What we will see


1. Qemu
  - a. Introduction
  - b. Developing a new simple emulated device
  - c. Adding it to an existing machine
2. Test
  - a. Simple bare-metal cortex a9 project
  - b. Test and debug with qemu
3. Test with Linux
  - a. Create a minimal rootfs
4. Develop a test a linux driver
  - a. Developing device driver
  - b. Adding it to an existing board configuring the device tree
  - c. Test

# QEMU

the FAST! processor emulator

## What is QEMU?

QEMU is a generic and open source machine emulator and virtualizer.



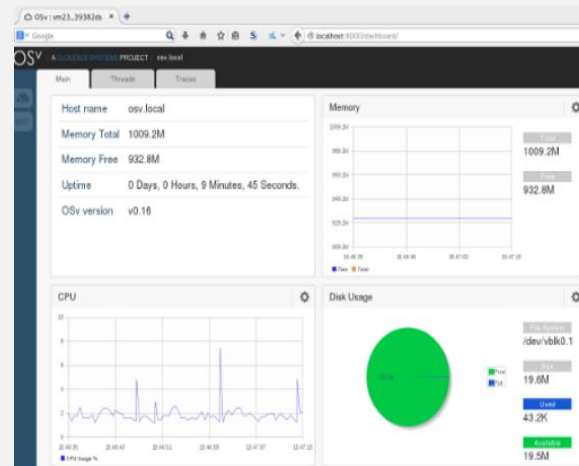
Full-system  
emulation

Run operating systems for any machine, on any supported architecture

```
[test@donizetti ~]$ qemu-arm ./ls --color /
bin etc lib64 mnt root srv system-upgrade-root var
boot home lost+found opt run sys
dev lib media proc sbin system-upgrade usr
[test@donizetti ~]$ uname -a
Linux donizetti 4.6.7-300.fc24.x86_64 #1 SMP Wed Aug 17 18:48:43 UTC 2016 x86_64
x86_64 x86_64 GNU/Linux
[test@donizetti ~]$ file ./ls
./ls: ELF 32-bit LSB executable, ARM, EABI5 version 1 (SYSV), dynamically linked
, interpreter /lib/ld-linux-armhf.so.3, for GNU/Linux 3.0.0, stripped
[test@donizetti ~]$
```

User-mode  
emulation

Run programs for another Linux/BSD target, on any supported architecture



Virtualization

Run KVM and Xen virtual machines with near native performance

# Device Specification

1. One memory mapped RW registry. You can read on it the last value write plus 1.
2. If you write 1 and interrupt is raised and writing 0 the interrupt is acknowledged

Device is installed in a custom board  
coming from Sambrelite board

This board is based on i.mx6 that is a ARM Cortex A9 MP processor so we will use a gcc for that platform for developing and testing

And now .... A brief introduction to QOM.

And jump to the C code

We have the device emulated and we can start to add it to a machine emulated.

So go back to the QEMU source code



# A bare metal CORTEX-A9

## startup.S

# And now LINUX

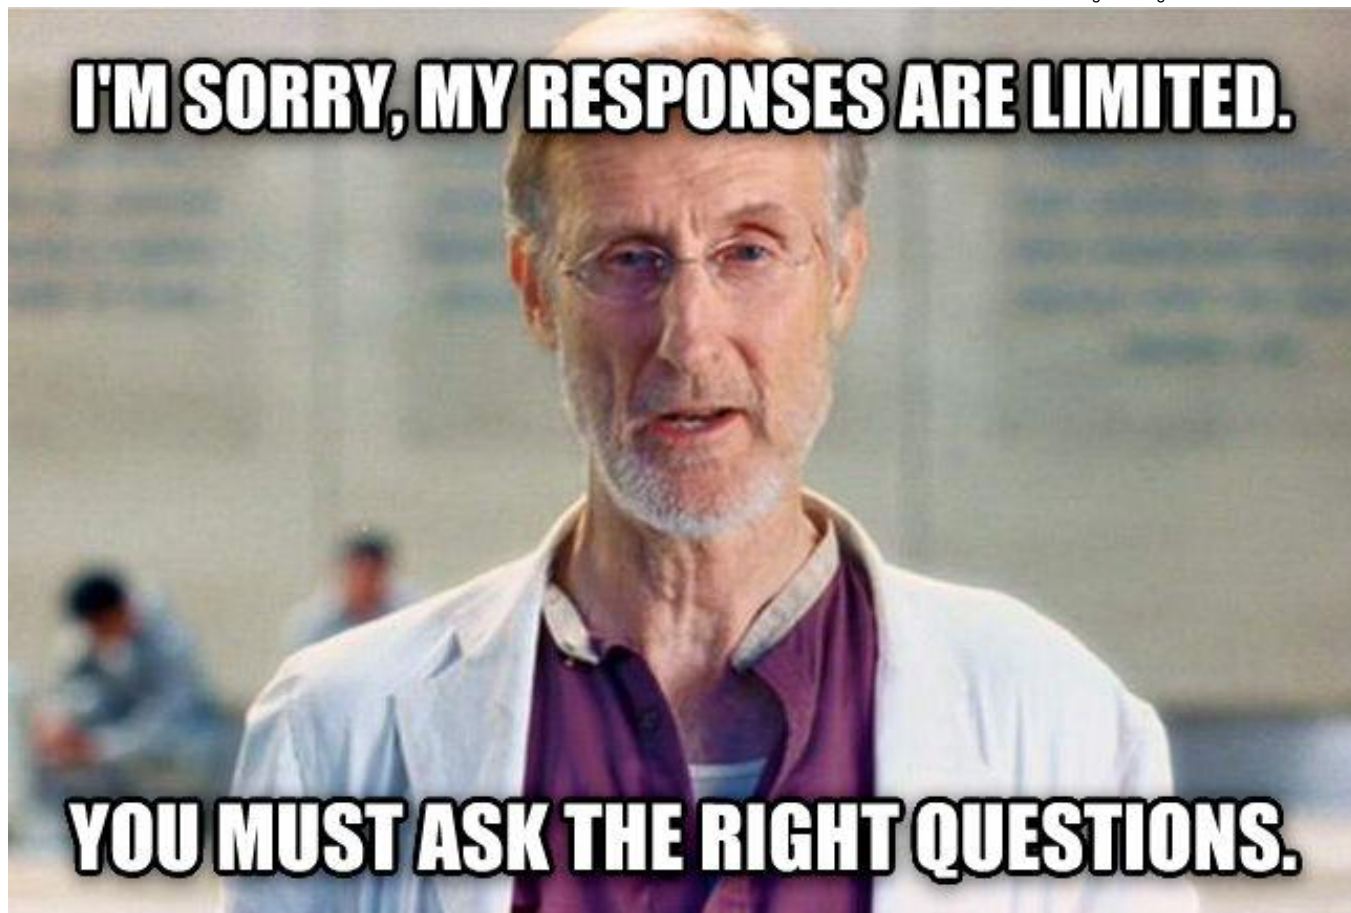
To develop and test a linux device driver we need a working Linux system.  
So before start to develop we have to create a minimal Linux system.

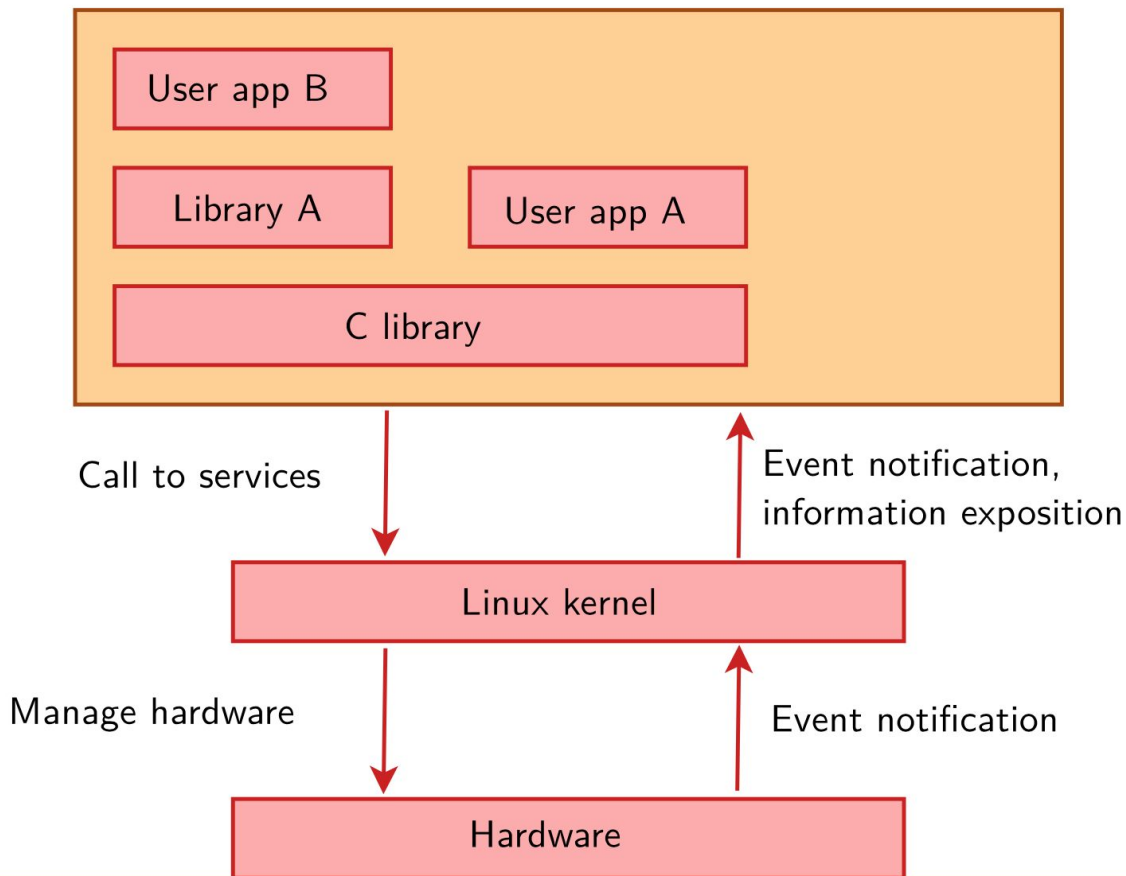
1. Build linux kernel
2. Build busybox
3. Build a minimal rootfs that can be mounted in ram. INITRAMFS

The result is not a full linux distribution but is something perfect for develop and test.

# Test the created Linux system

Questions ?






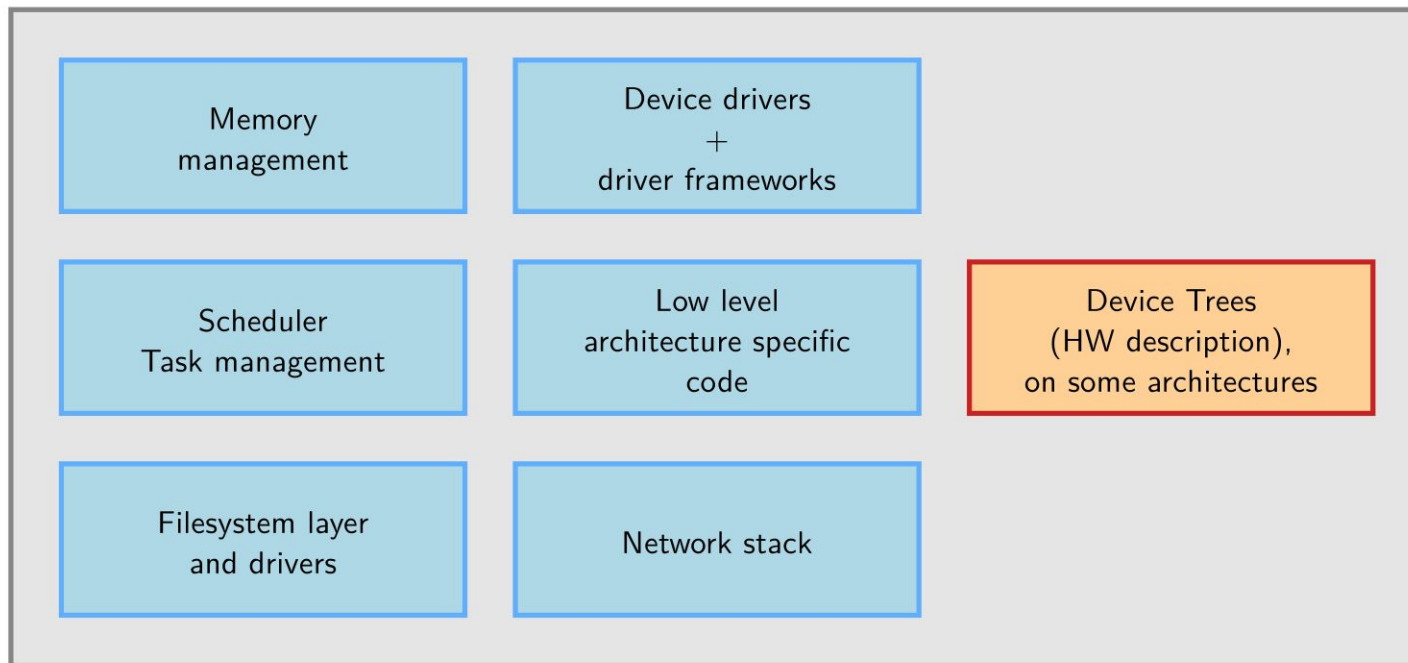
- ▶ Interact with the underlying hardware
- ▶ Provide a set of portable, architecture-independent API
- ▶ Handle concurrent access

- ▶ Kernel and userspace communicate using **system calls**, or *syscalls*
  - ▷ A modern Linux kernel implements around 300 syscalls
- ▶ Syscalls are propagated by the kernel to the dedicated **driver**
- ▶ Drivers handle the request and return a **status code** to the calling process

- ▶ Kernel can also present statistics and configurations using **pseudo-file systems**
- ▶ Usually mounted automatically at boot
- ▶ Several pseudo-file systems exist with different roles
  - ▷ `proc`, mounted on `/proc`
  - ▷ `sysfs`, mounted on `/sys`



## Linux Kernel



Implemented mainly in C,  
a little bit of assembly.



Written in a Device Tree  
specific language.

# Working on the Linux kernel

- ▶ By default, Linux is configured and built for the host system
- ▶ Use **ARCH** and **CROSS\_COMPILE** variables to change target platform
- ▶ Different configuration options are available for different platforms

- ▶ To use an existing configuration: **make <config-name>\_defconfig**
  - ▷ Configurations available in *arch/<arch>/configs/*
- ▶ To modify the current configuration: **make menuconfig**

- ▶ To build the kernel: **make -jN**
  
- ▶ Build outputs
  - ▷ `./vmlinux`
  - ▷ `./arch/<arch>/boot/*Image`
  - ▷ `./arch/<arch>/boot/dts/*.dtb`
  - ▷ `./**/*.ko`

# Writing a kernel driver

- ▶ (Optional) Create a dedicated directory under *drivers/*
- ▶ Create the necessary **Kconfig** options and **Makefile** rules
- ▶ Write code!

- ▶ A kernel module is not a userspace application!
  - ▷ Does not execute sequentially
  - ▷ Does no automatic cleanup
  - ▷ Can be interrupted
  - ▷ Does not have floating-point support



# Device drivers

- ▶ Access to the hardware is abstracted as a stream of bytes
  - ▷ Serial ports and text consoles are example of char devices
- ▶ Many hardware peripherals can be abstracted as char devices
  - ▷ Flash memories
  - ▷ Frame grabbers
  - ▷ ADC, DAC, shift registers

- ▶ Char devices are accessed using *device nodes*
  - ▷ Usually located in the `/dev` directory
- ▶ A combination of *major* and *minor* is used to identify the device
  - ▷ The major number identifies the device driver
  - ▷ Minors usually identify a device instance of a given driver

# Platform drivers

- ▶ A device located on a bus with minimal infrastructure
  - ▷ Most peripheral controllers on ARM SoCs are platform devices
- ▶ Usually addressed directly by the CPU
  - ▷ Eg. using memory-mapped registers
- ▶ In Linux, they are handled using **platform drivers**

- ▶ Implemented using **struct platform\_driver**
  - ▷ Registered using **platform\_driver\_register()**
- ▶ Enumeration is performed by the kernel
  - ▷ During enumeration, **.probe()** is invoked
  - ▷ During shutdown, **.remove()** is invoked
- ▶ The driver should actually check that the hardware *exists*

# Device Tree

- ▶ Embedded architectures often have non-discoverable peripherals
  - ▷ This is the case for ARM processors, for example
- ▶ We need a way to describe the hardware connected to such platforms
- ▶ Enter the Device Tree!



- ▶ Organized in `.dts` and `.dtsi` files
- ▶ Compiled in `.dtb` files
- ▶ Several roles in the Linux kernel
  - ▷ Describing devices attached to non-discoverable buses
  - ▷ Configure pin muxing on SoCs
  - ▷ Configure some system parameters (memory layout, flash partitioning etc.)

```
#include "imx6q.dtsi"
#include "imx6qdl-sabrelite.dtsi"

/ {
    model = "Freescale i.MX6 Quad SABRE Lite Board";
    compatible = "fsl,imx6q-sabrelite", "fsl,imx6q";

    soc {
        gemudev: qemu-test@50000000 {
            compatible = "linuxlab,qemu-test";
            reg = <0x50000000 4>;
            status = "okay";
        };
    };
};
```

# I/O operations

- ▶ I/O operations performed on device nodes are propagated to the driver
  - ▷ **open()**, **read()**, **write()**, **close()**, ...
- ▶ Drivers need to define a callback for each handled syscall
- ▶ This is done using the **struct file\_operations**
  - ▷ Structure members are all function pointers

- ▶ **open()** – called each time the device is opened from userspace
- ▶ **read()** – called when data is sent from the device to userspace
- ▶ **write()** – called when data is sent from userspace to the device
- ▶ **release()** – called when the device is closed in userspace

- ▶ Special care must be used to access user memory pointers in I/O callbacks
  - ▷ User memory is different from kernel memory, and can't be accessed directly
- ▶ **copy\_to\_user()** - Copy data from kernel memory to user memory
- ▶ **copy\_from\_user()** - Copy data from user memory to kernel memory

# Interrupt handling

- ▶ Most peripherals or devices use interrupts to signal asynchronous events
  - ▷ ADC conversion performed, SPI transfer completed, etc.
- ▶ In Linux, drivers can request to serve one or more specific interrupts
  - ▷ Usually each peripheral controller has its own interrupt line
- ▶ A device tree entry can be used to specify the interrupt index
  - ▷ Interrupt numbering is platform- and controller-specific